

November 23, 2023

# 1 Workbook

Use this notebook to complete the exercises throughout the workshop.

## Table of Contents

- [Section 1 – Getting Started with Pandas](#)
  - [Section 2 – Data Wrangling](#)
  - [Section 3 – Data Visualization](#)
- 

### 1.0.1 Section 1

#### Exercise 1.1

Create a DataFrame by reading in the `2019_Yellow_Taxi_Trip_Data.csv` file. Examine the first 5 rows.

```
[ ]: import pandas as pd
df = pd.read_csv("2019_Yellow_Taxi_Trip_Data.csv")
df[:5]
```

```
[ ]:  vendorid      tpep_pickup_datetime      tpep_dropoff_datetime \
0         2  2019-10-23T16:39:42.000  2019-10-23T17:14:10.000
1         1  2019-10-23T16:32:08.000  2019-10-23T16:45:26.000
2         2  2019-10-23T16:08:44.000  2019-10-23T16:21:11.000
3         2  2019-10-23T16:22:44.000  2019-10-23T16:43:26.000
4         2  2019-10-23T16:45:11.000  2019-10-23T16:58:49.000

   passenger_count  trip_distance  ratecodeid  store_and_fwd_flag \
0                1           7.93          1                N
1                1           2.00          1                N
2                1           1.36          1                N
3                1           1.00          1                N
4                1           1.96          1                N

   pulocationid  dolocationid  payment_type  fare_amount  extra  mta_tax \
0            138           170            1          29.5    1.0    0.5
1             11            26            1          10.5    1.0    0.5
```

2	163	162	1	9.5	1.0	0.5
3	170	163	1	13.0	1.0	0.5
4	163	236	1	10.5	1.0	0.5

	tip_amount	tolls_amount	improvement_surcharge	total_amount	\
0	7.98	6.12	0.3	47.90	
1	0.00	0.00	0.3	12.30	
2	2.00	0.00	0.3	15.80	
3	4.32	0.00	0.3	21.62	
4	0.50	0.00	0.3	15.30	

	congestion_surcharge
0	2.5
1	0.0
2	2.5
3	2.5
4	2.5

## Exercise 1.2

Find the dimensions (number of rows and number of columns) in the data.

```
[ ]: df.shape
```

```
[ ]: (10000, 18)
```

## Exercise 1.3

Using the data in the 2019\_Yellow\_Taxi\_Trip\_Data.csv file, calculate summary statistics for the fare\_amount, tip\_amount, tolls\_amount, and total\_amount columns.

```
[ ]: df[["fare_amount", "tip_amount", "tolls_amount", "total_amount"]].sum()
```

```
[ ]: fare_amount    151063.13
      tip_amount     26344.94
      tolls_amount    6234.47
      total_amount   225646.59
      dtype: float64
```

## Exercise 1.4

Isolate the fare\_amount, tip\_amount, tolls\_amount, and total\_amount for the longest trip by distance (trip\_distance).

```
[ ]: df.iloc[df["trip_distance"].idxmax(axis=0)][["fare_amount", "tip_amount",
↪ "tolls_amount", "total_amount", "trip_distance"]]
```

```
[ ]: fare_amount    176.0
      tip_amount     18.29
```

```
tolls_amount      6.12
total_amount      201.21
trip_distance      38.11
Name: 8338, dtype: object
```

---

## 1.0.2 Section 2

### Exercise 2.1

Read in the meteorite data from the `Meteorite_Landings.csv` file, rename the mass (g) column to mass, and drop all the latitude and longitude columns. Sort the result by mass in descending order.

```
[ ]: df2 = pd.read_csv("Meteorite_Landings.csv").rename(columns={"mass (g)":  
    ↳ "mass"}).drop(columns=["reclat", "reclong"]).sort_values(by=["mass"],  
    ↳ ascending=[False])  
df2
```

```
[ ]:
      name      id nametype      recclass      mass      fall \
16392      Hoba  11890   Valid      Iron, IVB  60000000.0  Found
5373      Cape York  5262   Valid      Iron, IIIAB  58200000.0  Found
5365  Campo del Cielo  5247   Valid      Iron, IAB-MG  50000000.0  Found
5370      Canyon Diablo  5257   Valid      Iron, IAB-MG  30000000.0  Found
3455      Armanty  2335   Valid      Iron, IIIE  28000000.0  Found
...      ...      ...      ...      ...      ...
38282  Wei-hui-fu (a)  24231   Valid      Iron      NaN  Found
38283  Wei-hui-fu (b)  24232   Valid      Iron      NaN  Found
38285      Weiyuan  24233   Valid  Mesosiderite      NaN  Found
41472  Yamato 792768  28117   Valid      CM2      NaN  Found
45698  Zapata County  30393   Valid      Iron      NaN  Found

      year      GeoLocation
16392  01/01/1920 12:00:00 AM  (-19.58333, 17.91667)
5373   01/01/1818 12:00:00 AM  (76.13333, -64.93333)
5365   12/22/1575 12:00:00 AM  (-27.46667, -60.58333)
5370   01/01/1891 12:00:00 AM  (35.05, -111.03333)
3455   01/01/1898 12:00:00 AM  (47.0, 88.0)
...      ...      ...
38282  01/01/1931 12:00:00 AM      NaN
38283  01/01/1931 12:00:00 AM      NaN
38285  01/01/1978 12:00:00 AM  (35.26667, 104.31667)
41472  01/01/1979 12:00:00 AM  (-71.5, 35.66667)
45698  01/01/1930 12:00:00 AM  (27.0, -99.0)
```

```
[45716 rows x 8 columns]
```

### Exercise 2.2

Using the meteorite data from the `Meteorite_Landings.csv` file, update the year column to only contain the year, convert it to a numeric data type, and create a new column indicating whether the meteorite was observed falling before 1970. Set the index to the id column and extract all the rows with IDs between 10,036 and 10,040 (inclusive) with `loc[]`. **Hint 1:** Use `year.str.slice()` to grab a substring.

**Hint 2:** Make sure to sort the index before using `loc[]` to select the range.

**Bonus:** There's a data entry error in the year column. Can you find it? (Don't spend too much time on this.)

```
[ ]: cleanedDf = df2.copy()
cleanedDf['year'] = cleanedDf['year'].apply(lambda x: int(x[6:11]) if
↳ isinstance(x, str) else x)
# cleanedDf['year'] = cleanedDf['year'].map(lambda x: x[6:11],
↳ na_action='ignore')
cleanedDf['before_1970'] = cleanedDf['year'].map(lambda y: y < 1970)
cleanedDf.set_index('id', inplace=True)

rows = cleanedDf.loc[(cleanedDf.index >= 10036) & (cleanedDf.index <= 10040)]
contains_non_string = df2.apply(lambda x: x.name if not isinstance(x['year'],
↳ str) else None, axis=1).dropna()
contains_non_string
rows
```

```
[ ]:
      name nametype      recclass    mass   fall   year  \
id
10039  Ensisheim   Valid          LL6  127000.0  Fell  1491.0
10038    Enshi     Valid           H5    8000.0  Fell  1974.0
10037    Enon     Valid  Iron, ungrouped    763.0  Found  1883.0
10036   Enigma     Valid           H4     94.0  Found  1967.0

      GeoLocation  before_1970
id
10039    (47.86667, 7.35)      True
10038    (30.3, 109.5)      False
10037    (39.86667, -83.95)      True
10036  (31.33333, -82.31667)      True
```

### Exercise 2.3

Using the meteorite data from the `Meteorite_Landings.csv` file, create a pivot table that shows both the number of meteorites and the 95th percentile of meteorite mass for those that were found versus observed falling per year from 2005 through 2009 (inclusive). **Hint:** Be sure to convert the year column to a number as we did in the previous exercise.

```
[ ]: 
```

```

dfConverted = df2.copy()
dfConverted['year'] = dfConverted['year'].apply(lambda x: x[6:11] if
↳ isinstance(x, str) else x)
dfConverted['year'] = pd.to_numeric(dfConverted['year'], errors='coerce')
dfConverted = dfConverted[(dfConverted['year'] >= 2005) & (dfConverted['year']
↳ <= 2009)]

dfConverted

pivot_table = pd.pivot_table(dfConverted,
                              values=['mass'],
                              index=['year'],
                              columns=['fall'],
                              aggfunc={'mass': ['count', ('95%-Quantil', lambda
↳ x: pd.Series.quantile(x, q=0.95))]}, fill_value=0)

# pivot_table.columns = ['fall_count_found', 'fall_count_observed',
↳ '95th_percentile_found', '95th_percentile_observed']

pivot_table

```

```

[ ]:
      mass
      95%-Quantil      count
fall      Fell      Found  Fell Found
year
2005.0      0.0  4500.00      0   874
2006.0  25008.0  1600.50      5  2450
2007.0  89675.0  1126.90      8  1181
2008.0 106000.0  2274.80      9   948
2009.0   8333.4  1397.25      5  1492

```

## Exercise 2.4

Using the meteorite data from the Meteorite\_Landings.csv file, compare summary statistics of the mass column for the meteorites that were found versus observed falling.

```

[ ]: dfCompare = df2.copy()
grouped = dfCompare.groupby('fall')['mass'].describe(include='all')
grouped

```

```

[ ]:
      count      mean      std  min  25%  50%  75%  \
fall
Fell  1075.0  47070.715023  717067.125826  0.1  686.00  2800.0  10450.0
Found  44510.0  12461.922983  571105.752311  0.0    6.94    30.5   178.0

      max

```

```

fall
Fell    23000000.0
Found   60000000.0

```

## Exercise 2.5

Using the taxi trip data in the `2019_Yellow_Taxi_Trip_Data.csv` file, resample the data to an hourly frequency based on the dropoff time. Calculate the total `trip_distance`, `fare_amount`, `tolls_amount`, and `tip_amount`, then find the 5 hours with the most tips.

```

[ ]: df4 = pd.read_csv("2019_Yellow_Taxi_Trip_Data.csv",
    ↳ parse_dates=['tpep_dropoff_datetime'])
df4 = df4.assign(hour=lambda x: x["tpep_dropoff_datetime"].dt.strftime('%H'))\
    .filter(["hour", "trip_distance", "fare_amount", "tolls_amount",
    ↳ "tip_amount"])

grouped = df4.groupby("hour").sum()
grouped.sort_values(by=["tip_amount"], ascending=False).filter(["tip_amount"]).
    ↳ head(5)

```

```

[ ]:      tip_amount
hour
16      12249.32
17      12044.03
18       1907.64
15         75.10
19         25.74

```

## 1.0.3 Section 3

### Exercise 3.1

Using the TSA traveler throughput data in the `tsa_melted_holiday_travel.csv` file, create box plots for traveler throughput for each year in the data. Hint: Pass `kind='box'` into the `plot()` method to generate box plots.

```

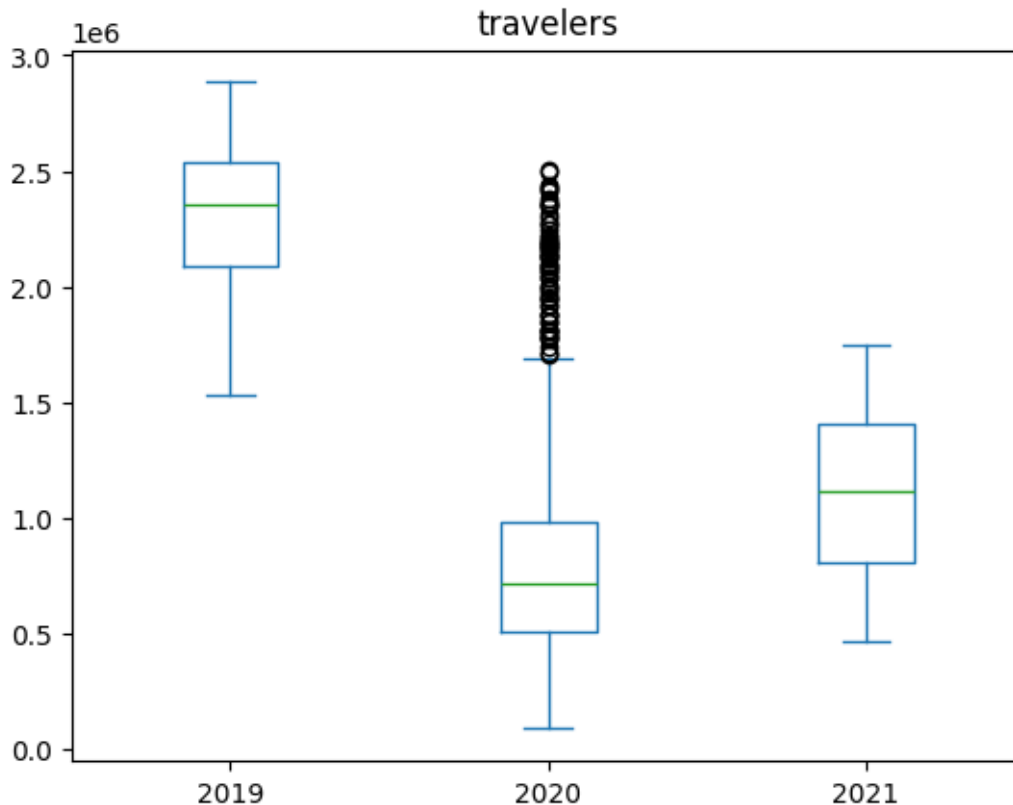
[ ]: dfTsa = pd.read_csv("tsa_melted_holiday_travel.csv", parse_dates=True)
dfTsa.plot(kind='box', column=["travelers"], by=["year"])
# dfTsa.boxplot(column=["travelers"], by=["year"], kind='box')

```

```

[ ]: travelers    Axes(0.125,0.11;0.775x0.77)
dtype: object

```



### Exercise 3.2

Using the TSA traveler throughput data in the `tsa_melted_holiday_travel.csv` file, create a heatmap that shows the 2019 TSA median traveler throughput by day of week and month.

```
[ ]: import seaborn as sns

working = pd.read_csv("tsa_melted_holiday_travel.csv", parse_dates=True)
# working.set_index('date', inplace=True) # unsere Idee, funktioniert aber nicht
working = pd.read_csv(
    'tsa_melted_holiday_travel.csv',
    parse_dates=True, index_col='date') # aus dem Internet, funktioniert.
    ↳ Warum bzw. wo der Unterschied zu uns liegt, ist nicht ganz klar.


year2019 = working.loc['2019']
year2019 = year2019.assign(week_day=lambda x: x.index.dayofweek + 1, month=
    ↳ lambda x: x.index.month)
year2019_pivot = year2019.pivot_table(
    values=['travelers'],
    index=['month'],
```

```

        columns=['week_day'], aggfunc='median'
    )
    sns.heatmap(data=year2019_pivot / 1e6,
                cmap='Blues',
                annot=True,
                xticklabels=["Mon.", "Tue.", "Wed.", "Thu.", "Fri.", "Sat.", "Sun."],
                yticklabels=["Jan.", "Feb.", "March", "May", "June", "July", "Aug.", "Sep.", "Oct.", "Nov.", "Dec."])

```

```
[ ]: <Axes: xlabel='None-week_day', ylabel='month'>
```



### Exercise 3.3

Annotate the medians in the box plot from *Exercise 3.1*. Hint: The x coordinates will be 1, 2, and 3 for 2019, 2020, and 2021, respectively. Alternatively, to avoid hardcoding values, you can use the `Axes.get_xticklabels()` method, in which case you should look at the [documentation](#) for the Text class.

```

[ ]: dfTsa = pd.read_csv("tsa_melted_holiday_travel.csv", parse_dates=True)
     plot = dfTsa.boxplot(column=["travelers"], by=["year"], return_type=None)

```



```
type(plot)
# plot.get_xticklabels()
```

```
[ ]: matplotlib.axes._axes.Axes
```

